FINAL PROJECT



Predicting Insurance Premium with given risk factors:
Pre-processing and Exploratory data analysis, Feature
Engineering, Feature Selection, Model, Pickle, Docker,
Flask

COURSE: INFO7390

Advance Data Science & Architecture

PROFESSOR:

Srikanth Krishnamurthy

SUBMITTED BY: TEAM 9

Amit Pingale - 001898697 Himani Solanki - 001899580 Shubham Patel - 001899476

Objective:

The Report summarizes the design and implementation of machine learning performed on the Insurance dataset. The documentation is divided into 6 parts:

Part 1: Feature Engineering

Part 2: Exploratory Data Analysis

Part 3: Regression Models

Part 4: Pickling the Model

Part 5: Pipeline

Part 6: Flask

Part 1: Feature Engineering

Step 1:Missing value analysis

Checking for missing values

Dataset has no missing values

Step 2:Converting necessary column to binary value for evaluation

```
In [40]: def smoking_habits(column):
              mapped=[]
               for row in column:
                   if row=="yes":
                       mapped.append(1)
                   else:
                       mapped.append(0)
               return mapped
          df["smoker"]=smoking_habits(df["smoker"])
In [41]: df.head()
Out[41]:
                           bmi children smoker
              age
                     sex
                                                   region
                                                              charges
           0
               19 female 27.900
                                              1 southwest 16884.92400
           1
               18
                    male 33.770
                                              0 southeast
                                                          1725.55230
               28
                    male 33.000
                                                 southeast
                                                           4449,46200
                    male 22.705
           3
               33
                                      0
                                              0 northwest 21984.47061
               32
                    male 28.880
                                                 northwest
                                                           3866.85520
```

```
In [43]: def gender(column):
               mapped=[]
               for row in column:
                   if row=="male":
                        mapped.append(1)
                   else:
                        mapped.append(0)
               return mapped
          df["sex"]=gender(df["sex"])
In [44]: df.head()
Out[44]:
                         bmi children smoker
              age
                  sex
                                                 region
                                                            charges
           0
              19
                    0 27.900
                                    0
                                            1 southwest 16884.92400
           1
               18
                    1 33.770
                                    1
                                            0
                                               southeast
                                                         1725.55230
               28
                     1 33.000
                                    3
                                               southeast
                                                         4449.46200
           3
               33
                       22.705
                                               northwest 21984.47061
               32
                     1 28.880
                                    0
                                               northwest
                                                         3866.85520
```

Step 3: Assigning region indexing

```
In [45]: def region(column):
    mapped=[]

    for row in column:

        if row=="northeast":
            mapped.append(1)
        elif row=="northwest":
            mapped.append(2)
        elif row=="southeast":
            mapped.append(3)
        else:
            mapped.append(4)

        return mapped
df["region"]=region(df["region"])
df.head()
```

Out[45]:		age	sex	bmi	children	smoker	region	charges
	0	19	0	27.900	0	1	4	16884.92400
	1	18	1	33.770	1	0	3	1725.55230
	2	28	1	33.000	3	0	3	4449.46200
	3	33	1	22.705	0	0	2	21984.47061
	4	32	1	28 880	0	0	2	3866 85520

Part 2: Exploratory Data Analysis

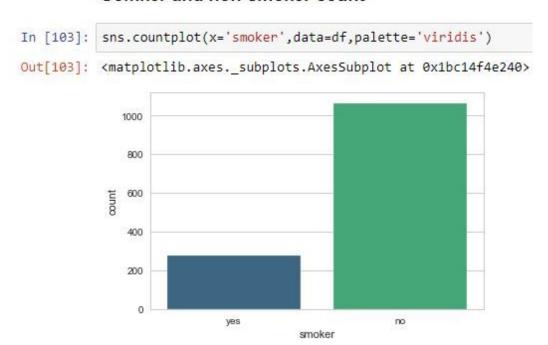
Step 1:Data info()

Step 2:Data describe()

In [101]:	df.des	cribe()			
Out[101]:		age	bmi	children	charges
	count	1338.000000	1338.000000	1338.000000	1338.000000
	mean	39.207025	30.663397	1.094918	13270.422265
	std	14.049960	6.098187	1.205493	12110.011237
	min	18.000000	15.960000	0.000000	1121.873900
	25%	27.000000	26.296250	0.000000	4740.287150
	50%	39.000000	30.400000	1.000000	9382.033000
	75%	51.000000	34.693750	2.000000	16639.912515
	max	64.000000	53.130000	5.000000	63770.428010

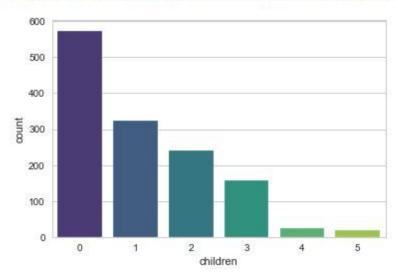
Step 3:

Somker and non-smoker count

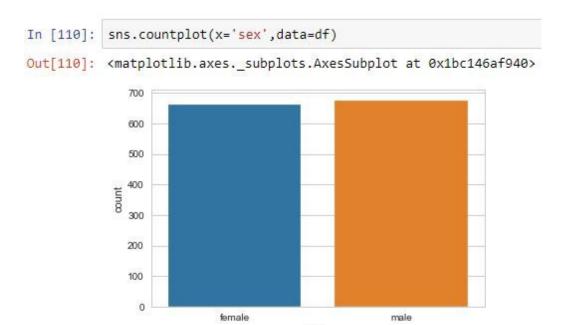


Children count

In [106]: sns.countplot(x='children',data=df,palette='viridis')
Out[106]: <matplotlib.axes._subplots.AxesSubplot at 0x1bc14eafe10>



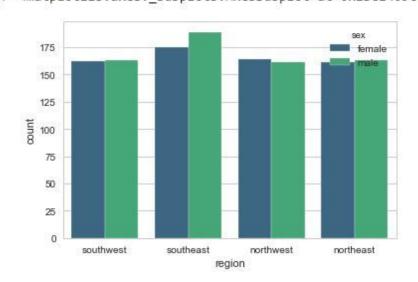
Step 4: Gender Count



Gender Count according to Region

```
In [111]: sns.countplot(x='region',data=df,hue='sex',palette='viridis')
Out[111]: <matplotlib.axes._subplots.AxesSubplot at 0x1bc1409ce80>
```

sex

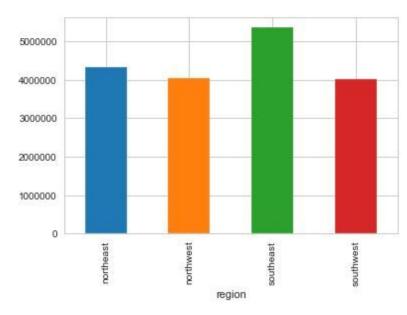


Step 5:

Health Insurance Charges according to Region



Out[112]: <matplotlib.axes._subplots.AxesSubplot at 0x1bc146af828>



Part 3: Regression Model

Step 1:Random Forest Regressor

```
In [22]: from sklearn.ensemble import RandomForestRegressor
         rf = RandomForestRegressor(n estimators = 1000, random state = 42)
         rf.fit(train features, train labels);
         predictions = rf.predict(test_features)
         # Calculate the absolute errors
         errors = abs(predictions - test labels)
         # Print out the mean absolute error (mae)
         print('Mean Absolute Error:', round(np.mean(errors), 2), 'degrees.')
         # Calculate mean absolute percentage error (MAPE)
         mape = 100 * (errors / test_labels)
         # Calculate and display accuracy
         accuracy = 100 - np.mean(mape)
         print('Accuracy:', round(accuracy, 2), '%.')
         from sklearn import metrics
         print('MAE:', metrics.mean absolute error(test labels, predictions))
         print('MSE:', metrics.mean_squared_error(test_labels, predictions))
         print('RMSE:', np.sqrt(metrics.mean_squared_error(test_labels, predictions)))
         print('R2:', metrics.r2 score(test labels, predictions))
         Mean Absolute Error: 2551.62 degrees.
         Accuracy: 73.36 %.
         MAE: 2551.615120945066
         MSE: 22499064.873433407
         RMSE: 4743.317918233334
         R2: 0.8508920741304946
```

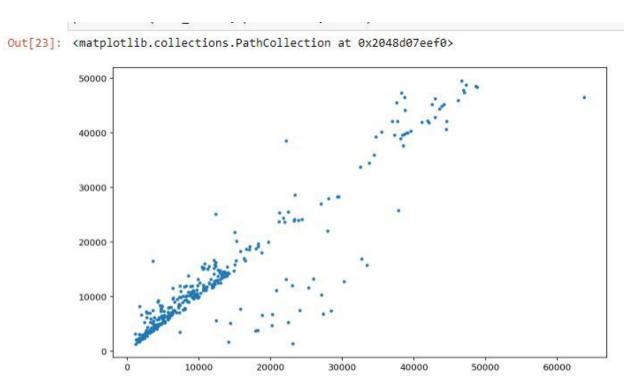


Fig. Scatter Plot for Random Forest Regressor

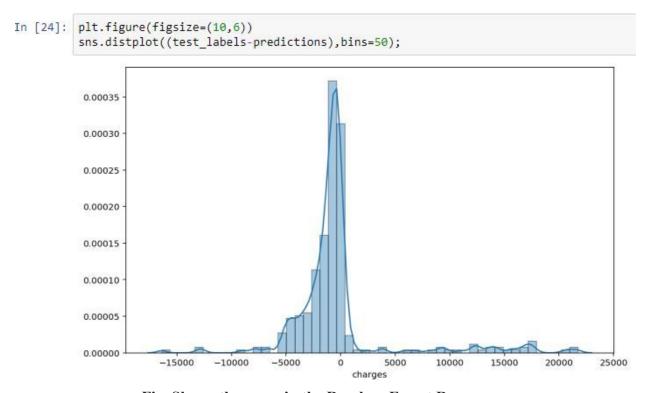


Fig. Shows the error in the Random Forest Regressor

Step 2: Neural Networks

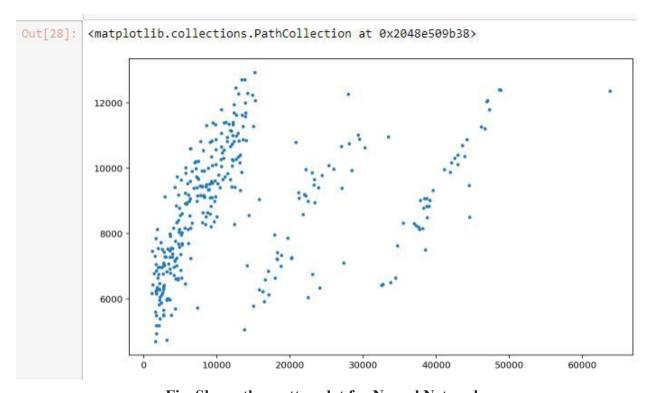


Fig. Shows the scatter plot for Neural Network

```
In [29]: plt.figure(figsize=(10,6))
    sns.distplot((y_test - y_test_prediction),bins=50);
```

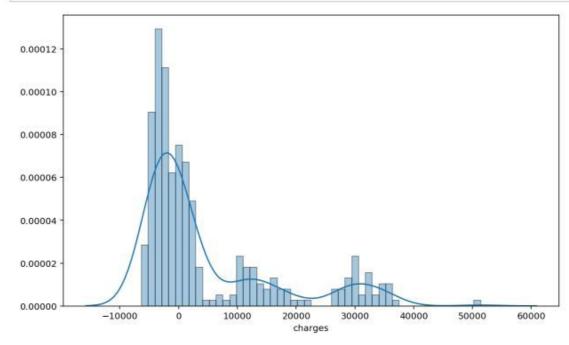
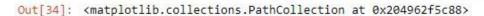


Fig. Shows the error for Neural Networks

Step 3:

Linear Regression

```
In [31]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
          from sklearn.linear_model import LinearRegression
          lm = LinearRegression()
          lm.fit(X_train,y_train)
          print(lm.intercept_)
          -11818.103976018108
In [32]: predictions = lm.predict(X test)
In [33]: from sklearn import metrics
          print('MAE:', metrics.mean_absolute_error(y_test, predictions))
          print('MSE:', metrics.mean_squared_error(y_test, predictions))
          print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
          print('R2:', metrics.r2_score(y_test, predictions))
         MAE: 4252.856455792365
         MSE: 35174149.327053055
         RMSE: 5930.779824530081
         R2: 0.7668905583460909
```



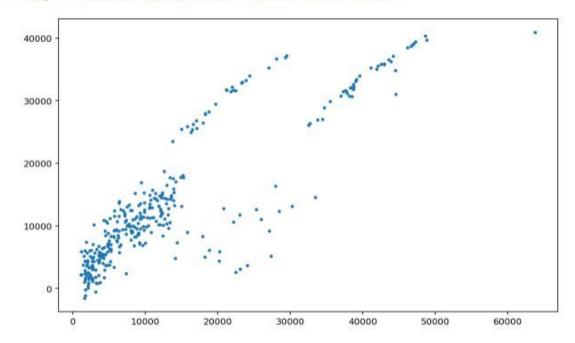


Fig. Showing scatter plot for linear regression

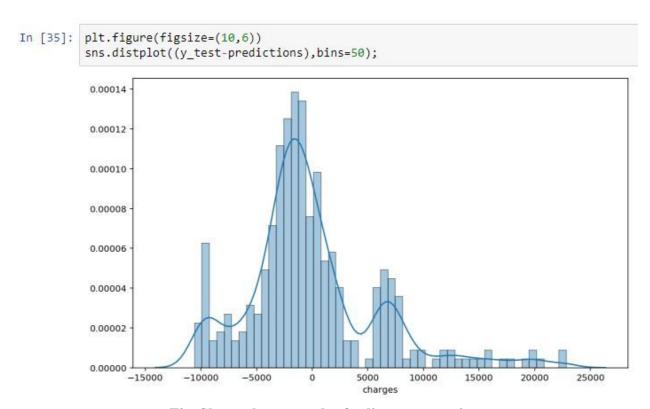
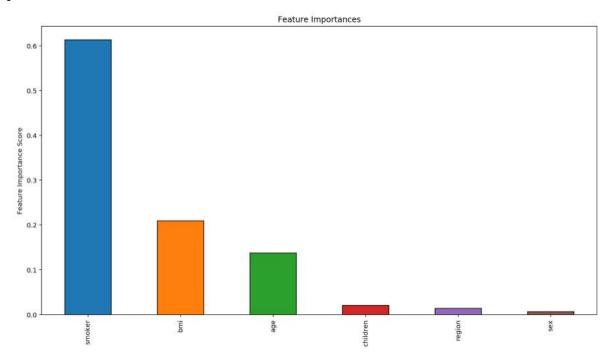


Fig. Shows the error plot for linear regression

<u>Step 4:</u>



Part 4: Pickle Model

Step 1:

Fitting Random Forest Model & Pickle for Clustor 0

```
from sklearn.ensemble import RandomForestRegressor
          dataset = pd.read_csv('Insurance_c0.csv')
In [6]:
           dataset.drop(["Unnamed: 0"], axis=1, inplace=True)
           dataset.head()
Out[6]:
               age sex
                             bmi children smoker region
                                                                  charges
                                         0
                                                  0
                                                           2 21984.47061
                33
                       1 22.705
            1
                32
                       1 28.880
                                         0
                                                  0
                                                               3866.85520
                       0 33.440
                                                               8240.58960
In [9]: rf = RandomForestRegressor(n_estimators = 100, random_state = 42)
         rf.fit(X_train,Y_train);
         filename = 'randforest c0.pckl'
         pickle.dump(rf,open(filename,'wb'))
         calc_error_metric('RandomForestRegression CO', rf, X_train, Y_train, X_test, Y_test)
         print('RandomForestRegression completed!')
        RandomForestRegression completed!
In [12]: from sklearn.linear model import LinearRegression
         linreg = LinearRegression()
         linreg.fit(X_train,Y_train)
         filename = 'linreg_c0.pckl'
         pickle.dump(linreg,open(filename,'wb'))
         calc_error_metric('Linear Regression', linreg, X_train, Y_train, X_test, Y_test)
         print('LinearRegression completed!')
         LinearRegression completed!
In [13]: from sklearn.ensemble import GradientBoostingRegressor
         gb = GradientBoostingRegressor(n_estimators=300,learning_rate= 0.1,max_features=1.0,random_state=42)
         gb.fit(X train, Y train);
         filename = 'gradboost_c0.pckl'
         pickle.dump(gb,open(filename,'wb'))
         calc_error_metric('Gradient Boosting Regression', rf, X_train, Y_train, X_test, Y_test)
        print('GradientBoostingRegression completed!')
```

GradientBoostingRegression completed!

Step 2:

Fitting Random Forest Model & Pickle for Clustor 1

```
In [14]: dataset = pd.read_csv('Insurance_c1.csv')
          dataset.drop(["Unnamed: 0"], axis=1, inplace=True)
          dataset.head()
 Out[14]:
             age sex bmi children smoker region
                                                  charges
                                            2 28923.13692
          0 60
                   0 25.84
          1
             62
                   0 26.29
                                            3 27808.72510
             56
                   0 39.82
                                0
                                            3 11090.71780
             52
                   0 30.78
                                      0
                                             1 10797.33620
          4 56
                   1 40.30
                                      0
                                             4 10602.38500
In [15]: X = dataset.drop(['charges'], axis =1)
          Y = dataset['charges']
In [16]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=np.random)
In [17]: rf = RandomForestRegressor(n_estimators = 100, random_state = 42)
          rf.fit(X_train,Y_train);
          filename = 'randforest_c1.pckl'
          pickle.dump(rf,open(filename,'wb'))
          calc_error_metric('RandomForestRegression C1', rf, X_train, Y_train, X_test, Y_test)
          print('RandomForestRegression completed!')
          RandomForestRegression completed!
In [18]: from sklearn.linear_model import LinearRegression
           linreg = LinearRegression()
           linreg.fit(X_train,Y_train)
           filename = 'linreg c1.pckl'
           pickle.dump(linreg,open(filename,'wb'))
           calc_error_metric('Linear Regression', linreg, X_train, Y_train, X_test, Y_test)
           print('LinearRegression completed!')
           LinearRegression completed!
In [19]: from sklearn.ensemble import GradientBoostingRegressor
        gb = GradientBoostingRegressor(n_estimators=300,learning_rate= 0.1,max_features=1.0,random_state=42)
        gb.fit(X_train,Y_train);
        filename = 'gradboost_c1.pckl'
        pickle.dump(gb,open(filename,'wb'))
        calc_error_metric('Gradient Boosting Regression', rf, X_train, Y_train, X_test, Y_test)
        print('GradientBoostingRegression completed!')
        GradientBoostingRegression completed!
```

Step 3:

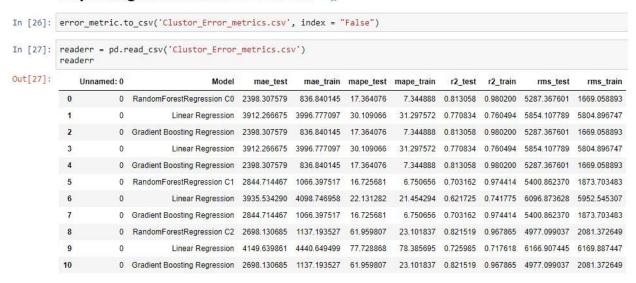
Fitting Random Forest Model & Pickle for Clustor 2

```
In [20]:
            dataset = pd.read_csv('Insurance_c2.csv')
            dataset.drop(["Unnamed: 0"], axis=1, inplace=True)
            dataset.head()
Out[20]:
                           bmi children smoker region
                age sex
                                                             charges
            0
                 19
                       0 27.90
                                       0
                                               1
                                                       4 16884.9240
                       1 33.77
                                                           1725.5523
             1
                 18
                                       1
                                               0
                 28
                       1 33.00
                                                           4449,4620
In [23]: rf = RandomForestRegressor(n_estimators = 100, random_state = 42)
          rf.fit(X train, Y train);
          filename = 'randforest_c2.pckl'
          pickle.dump(rf,open(filename,'wb'))
          calc_error_metric('RandomForestRegression C2', rf, X_train, Y_train, X_test, Y_test)
          print('RandomForestRegression completed!')
          RandomForestRegression completed!
In [24]: from sklearn.linear_model import LinearRegression
          linreg = LinearRegression()
          linreg.fit(X train, Y train)
          filename = 'linreg c2.pckl'
          pickle.dump(linreg,open(filename,'wb'))
          calc_error_metric('Linear Regression', linreg, X_train, Y_train, X_test, Y_test)
          print('LinearRegression completed!')
          LinearRegression completed!
In [25]: from sklearn.ensemble import GradientBoostingRegressor
        gb = GradientBoostingRegressor(n estimators=300,learning rate= 0.1,max features=1.0,random state=42)
        gb.fit(X_train,Y_train);
        filename = 'gradboost_c2.pckl'
        pickle.dump(gb,open(filename,'wb'))
        calc error metric('Gradient Boosting Regression', rf, X train, Y train, X test, Y test)
        print('GradientBoostingRegression completed!')
```

 ${\tt GradientBoostingRegression}\ {\tt completed!}$

Step 4:

Exporting model metrics csv file ¶



Part 5: Creating the pipeline

Step 1:Creating the pipeline

```
In [14]: | pipe_lr = Pipeline([('scl', StandardScaler()),('clf', LinearRegression(normalize=True))])
            grid params lr =[{}]
            gs lr = GridSearchCV(estimator=pipe lr, param grid=grid params lr, cv=10)
            gs_lr.fit(X_train, y_train)
            calc_error_metric('Regression', gs_lr, X_train, y_train, X_test, y_test)
            print('Regression completed')
           Regression completed
In [15]: pipe_rf = Pipeline([('scl', StandardScaler()),('rf', RandomForestRegressor(n_estimators=115,max_features=6,random_state=42))])
        grid_params_rf = [{}]
        gs_rf = GridSearchCV(estimator=pipe_rf, param_grid=grid_params_rf, cv=10)
        gs_rf.fit(X_train, y_train)
        calc_error_metric('RandomForest', gs_rf, X_train, y_train, X_test, y_test)
        print('RandomForest completed')
        RandomForest completed
grid_params_nn = [{}]
        gs_nn = GridSearchCV(estimator=pipe_nn, param_grid=grid_params_nn, cv=10)
        gs_nn.fit(X_train, y_train)
calc_error_metric('Nueral Network', gs_nn, X_train, y_train, X_test, y_test)
        print('Neural Network completed')
        C:\Users\shlok\anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:564: ConvergenceWarning: Stochastic
        Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
          % self.max_iter, ConvergenceWarning)
        C:\Users\shlok\anaconda3\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:564: ConvergenceWarning: Stochastic
        Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
```

Step 2: Calculate best model

```
In [18]: #### Calculate best model
         best model = min(rmse_dict.items(),key=operator.itemgetter(1))[0]
         print('Best Model is ', best model)
         print('Error Metrics are:')
         print(error metric)
         #### Write the error
         error_metric.to_csv('Error_metrics.csv')
         Best Model is RandomForest
         Error Metrics are:
                               Model
                                          mae test
                                                       mae train mape test
         0
                           Regression
                                       4238.237897
                                                     4170.916797
                                                                 46.846437
         0
                         RandomForest
                                       2338.356990
                                                     1075.872324 28.106875
                       Nueral Network 13430,703377 13063,905705 97,683536
         0
         0 GradientBoostingRegressor
                                       2560.283540
                                                    1723.767513 34.950779
            mape train
                        r2 test r2 train
                                               rms test
                                                            rms train
                                                          6061.345104
            41.274204 0.784400 0.740570
                                            5984.195122
         0
         0
             12.834150 0.892941 0.974542
                                          4216.904294
                                                          1898.753094
             97.798053 -1.085747 -1.204831 18612.835202 17670.418101
         0
             20.190387 0.882325 0.934318
                                           4421.029281
                                                          3049.873045
```

References:

https://www.nhlbi.nih.gov/health/educational/lose_wt/BMI/bmicalc.htm https://www.kaggle.com/mirichoi0218/insurance